

Auditing GBFS bike-sharing feeds at country and global scale: A reproducible anomaly taxonomy for open mobility data

Rohan Fossé^{*1} and Gaël Pallares²

¹CESI Engineering School, Montpellier, France

²CESI LINEACT (EA 7527), Montpellier, France

Abstract

Open mobility feeds are increasingly assumed to be directly ingestible into quantitative research pipelines: the General Bikeshare Feed Specification (GBFS) is now a regulatory requirement for French operators and is catalogued worldwide by MobilityData. We show that this syntactic standardisation is necessary but not sufficient. An audit of the 123 French GBFS systems combined with an exhaustive sweep of the MobilityData canonical catalogue (1,509 systems across 48 countries) yields a unified taxonomy of seven anomaly classes (A1–A7) covering out-of-domain inclusion, placeholder capacity, structural over-capacity, geospatial error, out-of-perimeter coverage, zero-capacity dock, and null capacity field. Across the French corpus the taxonomy reclassifies **30.9 %** of the raw stations (95 % bootstrap CI [30.5, 31.3]); across the global catalogue it flags 204 systems on A1–A5, 14 on A6 and 215 on A7, jointly covering 70,176 stations on A7 alone. The Italian deployment of a single operator illustrates the policy stake : 17 systems publish capacity = NaN on every station and would be “clean” under any v1.0 syntactic validator, yet collectively hide the most populous single-operator capacity blackout in the audited catalogue. We formalise a nine-step idempotent purging protocol that operationalises all seven anomaly classes and release the resulting *GBFS Audit Catalogue* (46,307 stations, 97 cities) on Zenodo. The Czech Republic is the European hotspot (55.6 % of audited systems flagged, all from a single operator, mirroring *Pony* in France). A positive control on Beryl London shows the standard can be used correctly when an audit tier exists. A single A3 correction collapses Bordeaux’s nominal docking total from 35,952 docks to 690 (−98.1 %), quantifying the policy cost of skipping the audit. The dataset, pipeline and Docker image are released for bit-exact reproduction.

Highlights

- Audit of 1,509 GBFS systems worldwide; unified taxonomy of 7 anomaly classes
- Hotspots are operator-driven: nextbike (Czechia), Pony (France)
- Single GBFS capacity field carries 4 incompatible semantics across 6 FR operators
- 30.9 % of raw FR stations removed after the nine-step purging protocol; further 61 % relabelled
- Audit Catalogue (46k stations) FAIR-released, code and Docker reproducible

Keywords: GBFS; data quality; open mobility data; FAIR; reproducibility; bike sharing; standards compliance.

^{*}Lead contact: rfosse@cesi.fr

1 Introduction

For more than a decade, bike-sharing systems have been a structural component of urban mobility policy in France and across Europe [1], advancing goals of liveability, decarbonisation and last-kilometre integration. To support inter-operability of the resulting fleets, the General Bikeshare Feed Specification (GBFS), introduced by the North American Bikeshare Association in 2015 and now maintained by the Open Mobility Foundation, provides a shared ontology covering station geolocation, dock counts, real-time availability and pricing [2, 3]. In France, the 2019 Mobility Orientation Law (LOM, article L. 1115-1 of the transport code) makes the publication of these feeds on *transport.data.gouv.fr* a regulatory acquis, so that by 2026 the working assumption of many research and policy pipelines is that GBFS feeds can be ingested directly.

That working assumption is misleading. The systematic audit reported here, covering 123 French GBFS systems, shows that standardised syntax masks substantial semantic heterogeneity. Three findings make the gap concrete. First, the same field (`station_information.capacity`) covers radically different realities: a physical dock count for dock-based fleets, a theoretical estimator obtained by conditional averaging on free-floating fleets, and arbitrary placeholders for under-reporting operators. Second, fourteen *Citiz* car-sharing systems are listed on the national portal under the GBFS schema, blurring the very definition of the object of study. Third, approximately 3.6% of the certified stations are flagged by A4 as 3- σ geospatial outliers from their system centroid, and four additional systems publish bounding boxes exceeding the 50,000 km² A5 threshold or list stations outside the national perimeter. Taken together, these observations support a single methodological claim: syntactic standardisation of an open feed is not sufficient to turn it into a research artefact, and any quantitative use of GBFS data must be preceded by an explicit, documented and reproducible audit.

Empirical audits of mobility feeds have so far remained piecemeal. Bike-sharing demand modelling and rebalancing analysis have generated a rich operational literature [4, 5, 6], but all of these works treat the input feeds as given and do not audit the standard itself. The closest precedent in transit-data quality is the smart-card literature : Pelletier et al. [7] systematically catalogue the gaps between what fare-collection feeds nominally publish and what they operationally support for research reuse, a programme structurally analogous to the GBFS audit we report here. Tao *et al.* [8] document sampling biases in mobile-phone records used for origin-destination estimation; Romanillos *et al.* [9] warn against statistical artefacts of ingesting raw GPS feeds directly; Bachand-Marleau et al. [10] surfaced operational-definition heterogeneity in bike-sharing demand data more than a decade ago; and Fishman [11] stresses the heterogeneity of operational definitions in international BSS comparisons. The North American Bikeshare and Scootershare Association tracks cross-operator harmonisation gaps annually [12], with the regulatory framing typical of an industry-side body but without the per-row audit trail a research consumer needs. In the adjacent transit literature, the GTFS community has been more proactive: the Canonical GTFS Schedule Validator [13] and the earlier Conveyal validator [14] codify hundreds of schema-level checks on transit feeds, and the data-engineering community has produced reusable declarative-validation frameworks such as Great Expectations [15], Pandera [16] and the Frictionless Data specifications [17]. The underlying data-quality literature is mature : Wang and Strong [18] established the consumer-centred dimensions, Pipino *et al.* [19] the assessment procedures, Khatri and Brown [20] the governance framing, Sebastian-Coleman [21] the ongoing-improvement perspective, and Naumann [22] the modern profiling approach. ISO/IEC 25012 [23] and its companion measurement standard ISO/IEC 25024 [24] formalise the quality dimensions ; for geographic information specifically, ISO 19157 [25] provides the corresponding framework, and the European INSPIRE directive [26] makes them regulatory. W3C's DCAT [27] and Data Quality Vocabulary [28] translate these dimensions into machine-readable metadata. None of these references, however, addresses GBFS specifically, and none produces a versioned, audited reference dataset at the scale of a country. The present paper fills

that gap for the French case and extends to the global GBFS catalogue.

The contributions are fivefold. First, an empirically-grounded **taxonomy** of seven anomaly classes (A1–A7) derived jointly from the audit of the 123 French GBFS systems and the 1,509-system MobilityData catalogue, and stress-tested on a ten-system live panel of non-French dock-based deployments (seven Spanish PBSC/Cyclocity networks plus Velo Antwerpen, Oslo Bysykkel and Bergen Bysykkel) which all pass the structural classes A1, A2, A3, A6 and A7 under unchanged thresholds. Second, a nine-step **purging protocol** (Algorithm 1) whose every transformation is logged, reproducible and reversible, and which jointly operationalises the seven classes. Third, the **Audit Catalogue GBFS France** data product (46,307 certified stations, 97 cities), released as a versioned Parquet file with DOI 10.5281/zenodo.20125460 under ODbL, accompanied by its JSON Schema, Croissant manifest and audit report, and mirrored as a Hugging Face dataset. Fourth, an **open-source pipeline** (`audit_pipeline` Python package, MIT, ~ 560 LOC, unit-tested) and a five-tab **interactive dashboard** (<https://gbfs-audit.streamlit.app>) that surfaces every audit verdict at the row level. Fifth, a **validation roadmap** of six falsifiable experiments, four of which already include a completed pilot, so that the protocol can be challenged in public rather than defended in private. Section 2 reports the empirical results, Section 3 discusses implications and limitations, and Section 4 details the experimental procedures.

2 Results

2.1 Anomaly taxonomy and incidence

The systematic audit of the 123 French GBFS systems uncovered five recurring anomaly classes, formalised in the `station_type` column of the Audit Catalogue and summarised in Table 1. Each class, in isolation, is sufficient to invalidate a non-audited cross-system comparison; the protocol that addresses them is given in Section 4.

Table 1: The seven GBFS anomaly classes (A1–A7) of the unified taxonomy, with their observed incidence on the released French catalogue (FR, 123 audited systems) and on the global MobilityData catalogue (Global, 1,509 systems across 48 countries). Each class is formalised in Methods (Section 4); per-system counts use the operational definitions documented in Section 4 (*e.g.* A3 flags any system with at least one `free_floating` entry after the S3 reclassification, which captures both natively-declared and post-audit-reclassified free-floating fleets). Counts are reproducible by re-running `python -m audit_pipeline.regenerate` on the released parquet.

Class	Name	Signature	FR	Global
A1	Out-of-domain inclusion	Car-sharing advertised as BSS	17	46
A2	Placeholder capacity	Constant non-zero c_i across docked subset	1	48
A3	Structural over-capacity	Conditional averaging on free-floating	41	33
A4	Geospatial error	$3\text{-}\sigma$ outlier from system centroid	78 (3.6% stns)	81
A5	Out of perimeter	System bounding box $> 50,000\text{ km}^2$	4	17
A6	Zero-capacity dock	$\geq 1\%$ of docked stations declare $c = 0$	0	14
A7	Null capacity field	$\geq 50\%$ of stations declare $c = \text{NaN}$	32 (FF)	215

The audit status of the 142 French GBFS feeds at the national portal is summarised in Figure 1: 126 systems are certified into the Audit Catalogue, 14 are excluded as micro-networks below the $N_{\min} = 20$ threshold, and 2 fail technical ingestion. The 14 car-sharing systems (A1) are filtered upstream.

A3 is the most pernicious anomaly. Free-floating fleets (*Pony*, *Cykleo*, etc.) advertise virtual stations in GBFS, the term “station” designating the GPS position of an anchoring point currently occupied by a vehicle. To avoid displaying empty stations, the operator typically reports a




Figure 1: Verdict of the systematic audit applied to the 142 GBFS feeds inventoried on *transport.data.gouv.fr* and *MobilityData*. 126 systems are certified; 14 are excluded as micro-networks (below the $N_{\min} = 20$ threshold) and 2 fail technical ingestion. The 14 car-sharing systems (A1) are filtered upstream of this catalogue.

capacity profile by *conditional averaging* on stations whose instantaneous capacity is non-zero:

$$\bar{c}_{\text{profile}} = \frac{\sum_{i: c_i > 0} c_i}{\#\{i : c_i > 0\}} \neq \bar{c}_{\text{actual}} = \frac{1}{N} \sum_{i=1}^N c_i. \quad (1)$$

Aggregated at the system level, $\bar{c}_{\text{profile}}$ wrongly classifies thousands of free-floating bikes as heavy dock-based stations. Cross-system comparison is therefore meaningful only after explicit reclassification; the empirical signature is visible in Figure 2 as a bimodal capacity distribution.




Figure 2: Empirical signature of the A3 floating-anchor bias. Declared capacity distribution by `station_type` (clipped at 80 docks) on the Audit Catalogue corpus. Only dock-based stations exhibit a physically interpretable capacity (median = 20).

2.2 Cumulative effect on the national corpus

Table 2 summarises the cumulative effect of the purging protocol. From 67,000 raw `station_information` entries published by the 142 inventoried French GBFS feeds, 20,693 (30.9%) are removed from the catalogue as falling outside the certification criteria (excluded micro-networks below $N_{\min}$, technical-ingestion failures, perimeter and outlier filters); the remaining 46,307 stations are kept

in the catalogue with an explicit `station_type` and per-row anomaly flags, of which a further 40,865 (61.0 % of the raw) are *relabelled* away from their declared GBFS type (39,235 relabelled to `free_floating` by step S3, 1,630 to `carsharing` by step S1). A 95 % cluster-bootstrap interval on the removal rate (resampling at the system level, $n = 10,000$, seed 42) is uninformative on this corpus because the 123 systems span four orders of magnitude in size (from 20 stations at the $N_{\min}$ threshold to $\sim 15,000$ for the largest Dott deployment); a single large system carries enough mass to flip the resampled mean. We therefore report the 30.9 % removal rate as a deterministic accounting fact on the audit-dated snapshot rather than as a population estimate. Re-running the audit on a freshly fetched snapshot is the right path to a temporal uncertainty assessment, which is the purpose of experiment E3 (twelve-month stability). The headline figure is that fewer than 9 % of raw entries survive into the catalogue without either being relabelled or removed — a strong empirical argument against the view that the standard alone delivers a research-ready dataset.

Table 2: Effect of the purging pipeline on the national GBFS corpus. *Certified* entries are kept in the Parquet with a typed `station_type` so that downstream consumers can filter explicitly; *rejected* entries are kept in a separate `rejected_stations.parquet` with the exclusion motive.

Indicator	Raw	Audit Catalogue
Inventoried GBFS systems	142	123 (certified)
Cities covered	—	97
with at least one dock-based station	—	64
Total stations (certified)	67,000	46,307
dock-based	—	5,442
free-floating (labelled, preserved)	—	39,235
car-sharing (A1)	1,630	re-labelled, excluded from BSS use
Stations reclassified or dropped	—	20,693 (30.9 %)

The most heavily impacted urban areas are those operated by free-floating fleets (*e.g.* Bordeaux, Nice, Saint-Étienne) and those historically lumped together with car-sharing systems on the national portal (Strasbourg, Grenoble), whereas urban areas served by traditional dock-based operators (Vélib’ in Paris, VéloMagg in Montpellier) are left almost intact. The certified-station mass by administrative region after purging is shown in Figure 3: Île-de-France concentrates the historical Vélib’ deployment, while regions with many free-floating systems (Nouvelle-Aquitaine, Auvergne-Rhône-Alpes) hold a much smaller dock-based stock after the A3 correction is applied.

2.3 Case study: Bordeaux and A3

Bordeaux is the paradigmatic case (Table 3). The raw feed assembles five operators across the agglomeration: *Pony* (free-floating), *Bird* (free-floating), *Dott* (free-floating), *Le Vélo par TBM* (dock-based) and *Citiz Gironde* (car-sharing). *Pony* alone advertises 2,996 `station_information` entries with declared capacities around 12 docks. Naive ingestion therefore credits Bordeaux with tens of thousands of nominal docking points and places it in the national top three for Vélib-class infrastructure.

The purging protocol, and in particular step S3 (the A3 correction), reclassifies the 2,996 *Pony* entries as `free_floating`, recomputes their actual capacity to 0.03 bike per entry, and leaves the 225 genuinely dock-based stations of the historical *Vélo par TBM* system. At the agglomeration level, this collapses the nominal docking total by 98 % and drops Bordeaux’s dock-based infrastructure from a top-quartile position on the national distribution of nominal docks to a position below the median: an order-of-magnitude effect on any supply-side aggregate built on the raw feed. The twelve urban areas most exposed to A3 are mapped in Figure 4.

Beyond the absolute station-count drop, the order-of-magnitude change in nominal infrastructure is the operational stake. Table 4 reports the five urban areas whose count of dock-based




Figure 3: Certified stations by administrative region, after the nine-step purging protocol. Annotations report the number of GBFS systems per region. Île-de-France concentrates Vélib’; Nouvelle-Aquitaine aggregates many systems but holds a smaller mass of dock-based stations once free-floating fleets have been re-typed by the A3 correction.

stations is most reduced by the purging protocol, together with the anomaly classes that drive the reclassification. The point is the consistency of the move: every one of these five urban areas would be credited with a top-decile dock-based deployment by a naive ingestion of GBFS, yet after the audit only Paris retains more than 1,500 certified dock-based stations. A non-audited pipeline does not just over-count by a small margin ; it changes the relative positioning of the principal French agglomerations on any supply-side metric by an entire quartile.

Two complementary case studies. A3 is not the only class with a measurable downstream effect. *Nice* carries the flagship A2 placeholder via `pony_Nice`: every station declares $c = 100$, yielding a total nominal capacity of about 25,000 bikes for an agglomeration whose actual dock-based network is essentially the rebadged *VéloBleu* of 91 stations. After purging, the nominal capacity collapses by a factor of 250 and Nice’s national rank moves out of the top quartile. *Strasbourg* provides the canonical A1 case: the *Citiz Grand Est* car-sharing system was historically ingested as part of the city’s BSS inventory, inflating the station count by roughly 35 %. After semantic exclusion, the agglomeration drops from rank 6 to rank 5, a smaller absolute move than Bordeaux but on a more discriminating part of the ranking. A1, A2 and A3 are therefore not interchangeable: each one distorts the ranking through a different mechanism (false positive on the population, on the capacity, or on the type), and a non-audited pipeline conflates them.

2.4 Comparison with naive baselines

A natural question is whether the full pipeline is justified or whether simpler heuristics would suffice. Three baselines are evaluated on the same 123-system corpus (Table 5). The *Raw GBFS* baseline ingests `station_information` as is. The *vehicle_type filter* mimics what a careful data engineer would write in a notebook. The *MobilityData audit* adapts the canonical GTFS-realtime schema-level rules to GBFS.

Three findings deserve emphasis. First, and most importantly, A3 reclassification recall is 0 % for every baseline (Wilson 95 % upper bound ≈ 8.6 % at $n = 41$ A3-positive systems): structural over-capacity is invisible to syntactic validators by construction. Only a semantically-

Table 3: Bordeaux: impact of the A3 correction on the *Pony Bordeaux* subsystem and on the agglomeration as a whole. The nominal docking total drops by 98.1 % once the *Pony* entries are correctly typed as `free_floating`; without the A3 step, a naive supply-side indicator would credit Bordeaux with about 36,000 physical docks, more than *Vélib’ Métropole* (the dock-based flagship of Île-de-France). The certified subset of 225 dock-based stations corresponds to the historical *Vélo par TBM* network.

Indicator	Raw GBFS	Audit Catalogue
<i>Pony Bordeaux subsystem</i>		
<code>station_information</code> entries	2,996	2,996 (relabelled)
Declared as dock-based	2,996	0
Declared capacity per entry (median)	12	12
Recomputed actual capacity	–	0.03 bike/entry
<i>Bordeaux agglomeration (all 5 systems)</i>		
Total raw entries	9,921	–
Certified dock-based stations	–	225
Nominal docking total (dock-based only)	35,952	690
Implied over-count factor	–	$\times 52$

Table 4: Top five French urban areas most affected by the purging protocol. “Raw” is the count of `station_information` entries in the unaudited feed ; “Cert. dock” is the count of dock-based stations in the Audit Catalogue. The drop is essentially the share of GBFS entries that the audit reclassifies as free-floating (A3), car-sharing (A1), or filtered (A4–A5). The “Flags driving the drop” column lists the anomaly classes that each agglomeration’s flagged systems trigger.

Urban area	Raw	Cert. dock	Drop	Flags driving the drop
Bordeaux	9,921	225	−97.7 %	A3 (Pony), A1 (Citiz), A7 (Dott, Bird)
Paris	19,872	1,507	−92.4 %	A7 (Dott, Bird), A3 (Pony Paris)
Marseille	2,391	221	−90.8 %	A3 (Pony), A7 (Dott)
Strasbourg	310	39	−87.4 %	A1 (Citiz Grand Est)
Lyon	2,057	452	−78.0 %	A3 (Pony, Bird), A7 (Dott)

aware step on the capacity distribution can recover it. Second, the *vehicle_type filter* and the *MobilityData audit* baselines produce identical numbers on this corpus, because schema-level checks do not catch anything that the name filter did not already exclude: French GBFS feeds are syntactically well-formed but semantically biased. Third, ρ_{FUB} is highest on the raw count and drops after audit; this is not a defect but a feature: FUB captures perception of the entire cycling environment, so the raw count over-fits the perceptual signal by including virtual stations, while the Audit Catalogue’s lower correlation reflects an honest count of physical docks.

2.5 Internal robustness checks

Three internal robustness checks are reported on the same corpus, ahead of the human inter-rater reliability that is the next priority of the validation roadmap.

Orthogonal-coverage. Three rule-based classifiers, each restricted to one signal source (name, capacity, geospatial), agree on a stratified 500-station sample at essentially chance level: pairwise Cohen’s κ near zero, Fleiss $\kappa = -0.07$, indistinguishable from independence. The sample is stratified jointly on `station_type` (`docked_bike`, `free_floating`, `carsharing`) and on the ten largest operators by station count, allocated proportionally to the size of each cell, with a floor of $n = 5$ per stratum to keep the smallest cells estimable; small operators are pooled into an `other` stratum. The seed used is `numpy.random.default_rng(42)` and the sample is shipped alongside the audit report so the agreement statistics can be independently recomputed. We




Figure 4: Top twelve French urban areas ranked by absolute reduction in dock-based station count between the raw GBFS feed and the certified Audit Catalogue. **Bordeaux** (highlighted) exhibits the largest absolute drop, driven by the reclassification of the *Pony* free-floating fleet (A3). Paris (Dott, Bird), Marseille (Pony, Dott) and Lyon (mixed) also see substantial corrections.

Table 5: Measured comparison of the Audit Catalogue against three baselines on the same 123-system French GBFS corpus. Brackets are 95 % confidence intervals (Wald on Dock, bootstrap on ρ_{FUB} over cities $n = 2,000$, Wilson on A1/A3). “Dock” for the Audit Catalogue is the reference by construction.

Pipeline	Dock	ρ_{FUB}	A1 ex.	A3 fix.
Raw GBFS	0.12 [.12, .12]	0.38 [.01, .68]	0 %	0 %
vehicle_type filter	0.15 [.15, .15]	0.19 [−.20, .52]	82 % [59, 94]	0 %
MobilityData audit	0.15 [.15, .15]	0.19 [−.20, .52]	82 % [59, 94]	0 %
Audit Catalogue	1.00	0.25 [−.13, .59]	100 % [82, 100]	100 % [91, 100]

do *not* interpret near-zero κ as a positive measure of “coverage” — inter-rater agreement is not a coverage statistic ; we report it only to falsify the alternative hypothesis that the three sources catch the same anomalies through correlated rules. The actually-relevant evidence is union recall: the three classifiers together reach 97.2 % recall at 91.5 % precision against the integrated pipeline, with only 7 of 500 stations flagged by all three. A1–A5 cannot be reduced to a single heuristic.

Full-task agreement. Three decision-tree classifiers seeing all signals but applying them in different priority orders reach Fleiss $\kappa = 0.80$ on the same sample (pairwise $\kappa \in [0.70, 1.00]$); disagreements concentrate on the A1/A3 frontier. The pair $\kappa \in [−0.07, 0.80]$ brackets the human inter-rater agreement that three independent annotators would plausibly achieve.

Threshold sensitivity. The A4 outlier detector uses a robust $3\text{-}\sigma$ rule on MAD-rescaled radii from the system centroid. We re-ran the audit under $\sigma \in \{2.0, 2.5, 3.0, 3.5, 4.0\}$ on the released parquet and report the result in Table 6. The absolute count of A4-flagged stations varies by a factor of three across the grid (3,182 at $\sigma = 2$ down to 1,084 at $\sigma = 4$), but the downstream

count-level summaries are essentially invariant : Kendall’s τ between the $\sigma = 3$ reference and each alternative ranges $[0.91, 1.00]$ on both the Top-10 cities and the Top-10 operators by certified docked-bike station count. The $\sigma = 3$ choice is therefore neither pivotal nor arbitrary; it sits on a plateau of stable downstream classification.

Table 6: A4 sensitivity sweep on the released v1.0.1 parquet (46,307 stations, 123 systems). τ_c and τ_o are the Kendall correlations of the Top-10 cities and Top-10 operators by certified docked-bike station count against the $\sigma = 3$ reference. Although the absolute A4 count varies threefold, the downstream ordering is invariant to within $\tau \geq 0.91$ across the entire grid.

σ	A4 stations	A4 systems	Share	τ_c	τ_o
2.0	3,182	100	6.87 %	0.911	0.911
2.5	2,185	91	4.72 %	0.911	0.911
3.0	1,655	78	3.57 %	1.000 (ref.)	1.000 (ref.)
3.5	1,305	62	2.82 %	0.964	0.964
4.0	1,084	55	2.34 %	1.000	0.956

These three checks do not replace human inter-rater reliability, which is the next priority of the validation roadmap (experiment E1, Section 3); they are reported as robustness measures, not as substitutes.

2.6 Cross-country transferability pilot

A natural objection to the present protocol is that A1–A7 might be artefacts of French publication practice rather than structural features of open-mobility-feed publication in general. To address this incrementally, a ten-system negative-control panel of European dock-based deployments is audited *live* via the public GBFS feeds, with the *same* rule set used on the French corpus and only the national perimeter recalibrated. All other thresholds ($\sigma = 3$, $N_{\min} = 20$, A2 placeholder rule, A3 capacity-profile ratio threshold of 5.0, $\tau_{A6} = 0.01$, $\tau_{A7} = 0.5$) are kept at their French values. The panel covers two large PBSC dock-based networks (Bicing Barcelona, bicimad Madrid), five smaller PBSC and JCDecaux Cyclocity deployments (Bilbao Bizi, Bizi Zaragoza, Dbizi Donostia, Sevici Seville, Valenbisi Valencia), the smartbike-powered Velo Antwerpen, plus the two original Norwegian pilots (Oslo and Bergen Bysykkel). Results are reported in Table 7; the live-audit script (`scripts/audit_live_systems.py`) and the per-system machine-readable results (`experiments/e5_europe/results.csv`) ship with the source repository, so the audit can be replayed against the same set of public feeds on demand.

Three observations follow. First, the capacity-profile ratio that defines the A3 signature on free-floating French fleets is essentially **unity** on both well-managed European dock-based systems, which empirically validates the discriminative power of Eq. 1. Second, vehicle-type metadata on both feeds is populated cleanly (Bicing: `{bicycle, scooter_standing}`; Oslo: `{bicycle}`), so no A1 false positive is triggered either. Third, perimeter and topological filters caught only a handful of edge outliers (1 station in Bicing, 3 in Oslo, 1 in Bergen), in line with what would be expected from any GPS-anchored fleet.

Empirical justification of $\tau_{A3} = 5$. The choice of the A3 threshold can be inspected directly on the released French corpus (Table 8). Of the 95 audited systems whose `station_information` populates `capacity`, 81 are dock-based and produce $\bar{c}_{\text{profile}}/\bar{c}_{\text{actual}} \in [1.00, 1.93]$ (median 1.00, 95th percentile 1.04). The remaining 14 are free-floating, of which 9 also sit near unity (capacity is either NaN or systematically zero, which makes the conditional and unconditional means coincide trivially) and 5 deploy the conditional-averaging publication pattern, giving ratios in the range $[51, 514]$. The histogram is bimodal with a wide empty region between roughly 2 and 50, and any threshold in that region produces an identical classification. We adopt $\tau_{A3} = 5$ at

Table 7: E5 panel: applying the A1–A7 rule set to ten non-French dock-based systems with French thresholds and country-appropriate perimeter only. The original three-system pilot (Bicing, Oslo, Bergen) is extended with seven additional systems audited *live* via the public GBFS feeds on 2026-05-13 using `scripts/audit_live_systems.py`. *None of the ten systems triggers the structural classes A1, A2, A3, A6 or A7*. A4 fires only on isolated stations ($\leq 4.5\%$ per system) and A5 fires once, on Sevici, exclusively as a consequence of two sentinel coordinates (0,0) that inflate the bounding-box area (an A4 contamination of A5 rather than a genuine $> 50,000\text{ km}^2$ deployment). The contrast with the 22+ French systems flagged on the same rule set on the 123-system Audit Catalogue corpus is substantial.

System	Country	Stations	A1	A2	A3	A4 (%)	A5	A6	A7
Bicing Barcelona	ES	545	—	—	—	0.18 %	—	—	—
bicimad Madrid	ES	637	—	—	—	0.00 %	—	—	—
Bilbao Bizi	ES	48	—	—	—	0.00 %	—	—	—
Bizi Zaragoza	ES	276	—	—	—	0.36 %	—	—	—
Dbizi Donostia	ES	70	—	—	—	2.86 %	—	—	—
Sevici Seville	ES	247	—	—	—	4.45 %	✓*	—	—
Valenbisi Valencia	ES	279	—	—	—	0.72 %	—	—	—
Velo Antwerpen	BE	321	—	—	—	0.00 %	—	—	—
Oslo Bysyssel	NO	266	—	—	—	$\leq 1\%$	—	—	—
Bergen Bysyssel	NO	97	—	—	—	$\leq 1\%$	—	—	—
French corpus (ref.)	FR	46,307	17 sys.	1 sys.	41 sys.	78 sys.	4 sys.	0 sys.	32 sys.

* Sevici’s A5 firing is a contamination artefact: two stations carry placeholder coordinates ($\text{lat} = 0, \text{lon} = 0$), which expand the system bounding box from $\sim 80\text{ km}^2$ (the true Seville network footprint) to $\sim 4,500,000\text{ km}^2$ (the great-circle span between Seville and the Atlantic origin). The A5 rule fires on the sentinel-inflated bbox; after manual removal of the two A4 outliers the A5 flag does not re-trigger. This empirically motivates the methodological recommendation in Section 3 that A5 be computed on A4-filtered coordinates rather than raw ones.

the lower edge of the gap to maximise robustness to small-corpus jitter without ever flagging a dock-based system on the released parquet.

Table 8: Empirical distribution of $\bar{c}_{\text{profile}}/\bar{c}_{\text{actual}}$ across the 95 audited French systems with non-empty **capacity**. The threshold band $\tau_{A3} \in [2, 50]$ partitions docked from biased free-floating systems with a margin of more than one order of magnitude on either side ; $\tau_{A3} = 5$ is reported in the paper for parsimony.

Ratio bucket	Docked systems	Free-floating systems
[1.00, 1.10]	79	9
[1.10, 2.00]	2	0
[2.00, 50.00]	0	0
[50.00, ∞)	0	5
Total	81	14

To turn the negative-control argument into a positive one, the rule set was then applied to the entire MobilityData canonical catalogue: **1,509 GBFS systems across 48 countries (the complete world inventory at the audit date)**. 1,421 systems (94.2 %) were reachable, 917 publish a `station_information` document, and **204 systems trigger at least one A1–A5 class** after three methodological corrections that emerged from the audit itself: a country-perimeter calibration on A4, a substring-matching fix in the A1 detector (the GBFS form factor `cargo_bicycle` contains the string `car` and was initially flagged as A1; the fix tightens A1 to exact list membership on the GBFS v3 form factor enum), and a convex-hull replacement of the bounding-box surface estimator in A5 (the rectangular bbox overestimates the operational

area for elongated networks; Switzerland’s national operators had a bbox area of 99,000 km² versus a hull area of $\approx 41,000$ km²). A side observation is that the French national portal *transport.data.gouv.fr* indexes only 123 of the 255 French entries that the MobilityData catalogue lists, so the global audit also exposes the incompleteness of the national regulatory pipeline. The corrected breakdown is shown in Table 9.

A3 (structural over-capacity) is detected on 29 non-French systems concentrated in Czech Republic (12) and Germany (9); A2 (placeholder capacity) on 34 non-French systems with Czech Republic again the hotspot (16); A1 (out-of-domain car-sharing labelled as BSS) on 29 non-French systems concentrated in the DACH region (Germany 21, Switzerland 8). Operator-level anti-patterns also travel internationally: *Dott* publishes `capacity = NaN` on every non-French deployment observed (Austria, UAE; 18 systems), and *Bird* does the same on three Canadian deployments. The cross-country audit also surfaces **A6 (zero-capacity dock)**, the sixth class of the taxonomy, formalised in Section 4.5. The empirical conclusion is that the documented anti-patterns are not French-specific but global, and the protocol generalises to a 1,254-system corpus with no methodological recalibration.

Table 9: Exhaustive audit of the seven-class taxonomy (A1–A7) on the MobilityData canonical catalogue (1,509 systems, including the 255 French entries). The per-country breakdown is shown for the four structural classes A1–A4 ; the macro-region class A5 (10 systems globally), the zero-capacity class A6 (14 systems) and the null-capacity class A7 (215 systems covering 70,176 stations) are calibrated on station-count-conditioned sub-corpora and are reported separately in Tables 13, 14 and in the country case studies of Section 2.6. Per-class counts are not disjoint (one system can trigger several classes) ; the “Flagged” column aggregates A1–A5 only. A4 is the geospatial out-of-perimeter class ; it requires more than 5 % of stations and at least 5 absolute outside-country stations to flag, to avoid edge-of-coast false positives. Countries with at least 30 audited systems or notable structural anomaly concentrations are shown.

Country	Audited	Reach.	Flagged	A1	A2	A3	A4 geo
France (FR)	255	253	38	17	14	4	3
Germany (DE)	251	243	32	21	4	9	–
United States (US)	175	172	0	–	–	–	–
Poland (PL)	94	91	28	–	3	–	25
Norway (NO)	83	46	2	–	–	–	2
Spain (ES)	69	66	1	–	–	–	1
United Kingdom (GB)	56	54	3	–	1	–	2
Finland (FI)	55	55	14	–	–	–	14
Netherlands (NL)	52	52	1	–	–	–	1
Switzerland (CH)	49	46	15	8	3	3	1
Czech Rep. (CZ)	45	43	25	–	16	12	–
Belgium (BE)	37	34	4	–	–	2	2
Italy (IT)	34	33	0	–	–	–	–
Austria (AT)	27	23	1	–	1	1	–
Sweden (SE)	23	23	12	–	–	–	12
Slovakia (SK)	19	19	16	–	–	–	16
Total (48 countries)	1,509	1,421	204	46	48	33	81

Czech Republic case study: one operator drives an entire country. The Czech Republic is the most striking national hotspot of the global audit, with 25 of the 45 audited systems flagged (55.6 %). The structural finding is that **all 45 audited Czech systems belong to a single operator**, *nextbike*, and that 25 of those deployments simultaneously trigger A2 (placeholder capacity, 16 systems) or A3 (structural over-capacity, 12 systems), with 3 deployments flagged by both classes. The 25 flagged Czech systems together publish 1,606 GBFS station entries that,




Figure 5: Per-country incidence of the structural classes A1–A4 across the 1,509-system MobilityData catalogue (top 18 countries by total flagged). The chart shows the stacked composition of the **Flagged** column of Table 9 and makes the country-level signatures legible at a glance: A4 (geospatial) drives the Polish, Finnish, Swedish and Slovak hotspots (cross-border bounding boxes on aggregator feeds), A1 (car-sharing labelled as BSS) drives the DACH cluster (Germany, Switzerland), A2 (placeholder capacity) and A3 (over-capacity) drive France and the Czech Republic. Macro-region (A5), zero-capacity (A6) and null-capacity (A7) are reported separately because they are calibrated on station-count-conditioned sub-corpora (Tables 13 and 14).

under the declared capacity semantics of the standard, would naively be counted as physical docks. This is the cross-country mirror of the French case: *Pony* drives the French free-floating A2/A3 hotspot, and *nextbike* drives the Czech one. The operator-specific structure of the hotspots is empirically consistent with the broader claim that the documented errors propagate through corporate publication conventions rather than through national regulatory practice.

Switzerland: heterogeneous fragmentation as the dual of operator monopoly. Switzerland is the only country in the audit, with France, for which all five A1–A5 classes are simultaneously detected: 8 A1, 3 A2, 3 A3, 1 A4 and 5 A5 across 16 flagged systems out of 49. The underlying pattern is the dual of the Czech case. Where Czech Republic has 45 systems published by a single operator that propagates one anti-pattern (*nextbike*, all A2/A3), Switzerland has **14 distinct entities publishing 16 flagged GBFS feeds**, two of which are federal multi-operator aggregators rather than single fleets. We distinguish three publisher categories. First, national-scale car-sharing brands (*2EM Car Sharing* 1,741 stations, *Mobility* 1,691, *MyBuri* 2,118, *edrive carsharing*, *Car-ship*, *QuickRent*) are listed on the same federal GBFS catalogue as bike-sharing fleets and trigger A1. Second, mixed-fleet operators (*Voi Switzerland* 771 e-scooter anchors, *Bolt Switzerland* 54, *nextbike Switzerland* 380) publish virtual-station entries with

capacity profiles that trigger A3. Third, *federal multi-operator aggregators* (*sharedmobility.ch* 12,091 entries, *Velospot* 1,634 entries) combine many operators into a single GBFS feed whose union-bounding-box necessarily exceeds the A5 macro-region threshold: A5 here is therefore a property of the *aggregation layer*, not of the underlying operators it republishes. The case is empirically informative because the same A1–A5 envelope can be reached through two opposite organisational configurations: monopoly publication (Czech Republic) or fragmented publication compounded by federal aggregation (Switzerland). The audit framework is publisher-agnostic and detects both, and the operator vs aggregator distinction is recorded in the audit report so downstream consumers can disambiguate the A5 attribution.

Italy: a clean audit that the capacity field cannot speak for. Italy is the most informative *boundary case* of the global audit: 33 of 34 systems reachable, 25 with published station data, and zero classed as flagged by A1–A6. We emphasise that this result is *not* a violation of the GBFS specification: every Italian system audited is spec-compliant. A closer look at the underlying capacity column nevertheless reveals a downstream research-reuse problem. Of the 25 Italian systems with data, 17 are deployments of *Dott* (Milan alone publishes 11,199 stations) and every *Dott* deployment declares `capacity = NaN` on 100 % of its `station_information` entries. Because A2 requires a constant non-zero value, A3 requires numerical capacities to compute the profile ratio, A4–A5 are geospatial, and A6 tests $c = 0$ (not $c = \text{NaN}$), the NaN-capacity pattern is not actionable from A1–A6. A7 fills that gap as a consumer-side warning (Section 4.6): **215 reachable systems with at least 20 stations cross the 50 % NaN-capacity threshold and are flagged exclusively by A7**, together covering 70,176 stations of which $\sim 46,000$ ($\sim 65\%$) are a single operator (*Dott*, 141 systems across Germany, Italy, Austria, UAE, and beyond) and the remaining $\sim 24,000$ are spread across 74 smaller systems. The point of A7 is not to label these systems non-compliant ; it is to make their capacity column non-aggregable to any downstream consumer who would otherwise sum NaN with non-NaN values silently.

Germany: how the audit caught itself. The German case carries a methodological lesson. An initial run of the global audit flagged 64 of the 251 German systems, of which 58 under A1 (out-of-domain car-sharing). Manual inspection of the flagged systems revealed that the A1 detector relied on substring matching ("`car`" in `form_factor`) and incorrectly fired on the GBFS v3 form factor `cargo_bicycle`, a cargo bike. After tightening A1 to exact list membership on the GBFS v3 form factor enum, the German count drops to 21 systems (Conficars, Stadtmobil, TeilAuto Greater Stuttgart, Ford Carsharing, Flinkster and similar brands) and the global flag count from 215 to 173. The episode illustrates two complementary points: open audit pipelines need their own audit (the rule of A1 was rule-level brittle), and the underlying signal is genuine (21 real car-sharing systems in Germany are still labelled as bike-sharing on national portals).

2.7 Operator-level fragmentation of the capacity field

Across the case studies above, a recurring mechanism explains why the same anomaly recurs across operators: the GBFS field `station_information.capacity` is populated by French free-floating operators according to four mutually incompatible publication conventions, observed across six operator–city combinations (Table 10), aggregating to **84.7 % of the certified French rows**. A downstream consumer who applies a single arithmetic to the field therefore aggregates four physically distinct quantities (NaN, constant placeholder, per-vehicle ratio, conditional fleet profile), silently mixed across six operator deployments. The corresponding operator-level pattern is global: Beryl London (UK) provides a positive control by publishing an empty `station_information` document for its free-floating fleet (zero entries), which is the canonical GBFS v3 behaviour for non-docked fleets and is therefore not flagged by any of A1–A7 (in particular A7 only fires on *populated* documents whose entries declare `capacity = NaN`, which

is the qualitatively different Dott pattern). The official MobilityData GBFS Validator [13], by contrast, returns no error on any of the flagged French systems (0 % recall on A3 in Table 5).

Table 10: Six observed semantics of the same `station_information.capacity` field across French free-floating operators. Total = 39,235 stations, 84.7 % of the certified corpus.

Operator	Cities	Stations	Reported $\bar{c}$	Semantics
Dott	8	20,224	NaN	Field left unfilled
Bird	11	4,499	NaN	Field left unfilled
Pony (Nice)	1	412	100	Constant placeholder (A2)
Pony (Paris)	1	3,617	1.6	Per-vehicle ratio
Pony (Bordeaux, Angers, ...)	12	7,436	2–15	Conditional fleet profile (A3)
Voi	4	2,816	5–10	Small fleet-profile estimator




Figure 6: Conceptual diagram of the four incompatible publication conventions of the GBFS `station_information.capacity` field observed across six French free-floating operator deployments (39,235 stations, 84.7 % of the certified rows). Each card pairs the operator, the convention applied to the field, the typical reported value, and the anomaly class triggered by that convention (A2 placeholder, A3 over-capacity profile, A7 null capacity). The diagram makes explicit why a downstream consumer applying a single arithmetic to `capacity` aggregates four physically distinct quantities silently mixed across operators; it complements Table 10 and motivates the dual-column `capacity_raw/capacity_audited` schema of the released catalogue.

3 Discussion

3.1 Implications

The audit is intended as a scientific validation tier between GBFS publication and academic use, addressing the gap between regulatory publication and research-grade reuse that the open-government-data literature identifies as a recurrent adoption barrier [29]. Three concrete avenues would institutionalise it: a jointly published *GBFS Audit Best Practices Handbook* by francophone micromobility laboratories to align practices currently rebuilt in isolation; a working group under ADEME or *transport.data.gouv.fr* to formalise an extended GBFS schema with a certification

framework; and a continuous validation chain running the pipeline monthly on the national feeds, with automated publication of the audit report and operator-side notification on regression. Star [30] reminds us that the most durable infrastructures are those whose governance is explicit and tied to an identifiable community of practice; moving the artefact from author-led to community-led governance is therefore an institutional challenge as much as a scientific one. The technical lesson of the audit, that a single field (`capacity`) cannot serve six different operational semantics, is also the empirical case for what Stonebraker and Çetintemel [31] call the end of “one size fits all” data architectures : open standards that attempt to span dock-based, free-floating and car-sharing fleets in the same schema accumulate semantic ambiguity exactly because they refuse to specialise.

3.2 What this paper claims — and what it does not

Three demarcations should be stated explicitly to avoid over-reading the contribution.

This is an inductive taxonomy, not a statistically validated one. A1 to A4 were derived from an exploratory scan of the 123 French systems; A5 from enrichment with the *transport.data.gouv.fr* catalogue; A6 and A7 from the empirical distributions surfaced by the global 1,509-system sweep (Section 4, Provenance of the taxonomy). Detection rules were calibrated on the same data they are reported to flag. The ten-system live cross-country panel (Section 2.6, Table 7) is an out-of-sample check on the negative-control side: every system in the panel was selected for its known publication conventions, then audited *live* via its public GBFS feed under the unmodified rule set, and every system passes A1, A2, A3, A6 and A7. But the construct-validity question still cuts deeper: the rules themselves were derived inductively, and the question is whether the same rules would survive on *positive* examples drawn from systems added to the catalogue after a rule freeze. The pre-registered protocol for that experiment (E1, including freeze hash, falsification conditions and a six-month timeline) ships in `experiments/e1_holdout/PROTOCOL.md` of the source repository ; the test will be executed against the next MobilityData snapshot and reported as a stand-alone appendix in the next revision of this manuscript. A reader who needs a strict separation between rule-derivation and rule-evaluation should treat the present catalogue as descriptive ground-truth rather than as the outcome of a held-out test, until the E1 results are available.

This is an audit of an open specification, not an audit of its operators. A2, A3, A6 and A7 detect patterns the GBFS specification *permits* (a constant placeholder, a conditionally-averaged virtual-station fleet, a zero-dock station, an empty capacity field). What our framework flags is that those permitted patterns are mutually incompatible with the arithmetic a downstream research consumer would naturally apply to the field. The technical lesson is about the standard, not about the publishers ; criticism of the operators we name is out of scope for this paper. The standards-side recommendation in Section 3 is therefore phrased as a request for specification revision, not for operator sanction.

This is a static audit. The `station_information` feed describes the deployment, not its dynamic behaviour. Real-time availability, rebalancing patterns and per-vehicle relocation live in the `station_status` feed, which carries its own quality issues and would require an analogous but disjoint audit protocol. The dynamic audit is the subject of the L2 limitation and the E4 future-work item.

3.3 Intellectual lineage in adjacent fields

The protocol’s design intent is consistent with two adjacent traditions of data audit that the mobility-data community has so far engaged with only sparingly.

Property-based testing in software engineering. Each of the seven anomaly classes A1–A7 admits a translation into a *property* that a compliant feed should satisfy: capacity is either a positive dock count, NaN or the same value across the system; coordinates lie within an

a-priori perimeter; the $\bar{c}_{\text{profile}}/\bar{c}_{\text{actual}}$ ratio is bounded. This is the same intellectual move as in QuickCheck-style property-based testing for software, where invariants on the output are stated explicitly and checked across many inputs [32]. The Great Expectations and Pandera toolchains [15, 16] we benchmarked against are descendants of the same lineage. The contribution of the present audit, in that light, is to identify which properties GBFS implicitly assumes and to make their absence diagnosable; the SE community has spent two decades building infrastructure for that, but the mobility-data community has not yet imported it.

Regulatory data audit in clinical-trial pipelines. The audit-trail / reversibility / per-step logging design of the purging protocol mirrors the audit requirements of clinical-trial data management standardised in ISO 14155 and in the FDA’s 21 CFR Part 11. The open-mobility-data community is, in effect, re-discovering with two decades of delay what the pharmacovigilance community formalised in the 1990s. The analogue is informative because it suggests that long-term sustainability of the audit protocol depends on a community of practice and a regulatory backstop (e.g. ADEME or the European EuroMobility coordination layer), not on the lifespan of any single research artefact.

3.4 When a clean verdict hides a missing column

The Italian case study of Section 2.6 crystallises what is at stake. Of the 34 Italian systems in the catalogue, 33 are reachable and 25 publish station data ; *none* of them is flagged by classes A1 to A6 in the original v1.0 taxonomy. A reviewer reading only the A1–A6 verdict would conclude that the Italian GBFS publication is the cleanest in the catalogue. The seventh class A7 reveals the opposite: 17 of the 25 data-publishing Italian systems are deployments of a single operator (*Dott*) whose Italian feeds declare `capacity = NaN` on 100 % of their $\sim 11,000$ `station_information` entries. The country is not clean ; the country is silent on the only field that matters for infrastructure quantification. An audit framework that conflates *no problem detected* with *no data published* mis-reads the entire national corpus. A7’s empirical role is precisely to keep those two cases distinguishable, and the Italian episode is the clearest reminder that absence of evidence is not evidence of absence in open-standards compliance auditing.

3.5 GBFS v3 and the durability of the v2 legacy

A natural counter-argument to the audit is that GBFS v3 (released 2023) introduces `is_virtual_station`, a boolean that in principle disambiguates free-floating anchors from physical docks and would mechanically suppress A3 detection. Three observations temper this expectation. First, **adoption is slow**: of the 1,509 audited systems, only 541 (35.9 %) declare GBFS v3 in their auto-discovery feed at the audit date, against 772 (51.2 %) on v2.x ; 50 remain on v1 and 146 expose no parseable version. Second, the v2 legacy will persist: regulatory feeds on *transport.data.gouv.fr* and the MobilityData canonical catalogue archive snapshots that any longitudinal study from 2017 onward will need to ingest, and a v3 feed published today does not retro-fix the v2 historical record. Third, and most importantly, **adoption is not compliance**: the German cargo-bicycle mis-detection we report (Section 2.6) and the Italian *Dott* NaN-capacity propagation both occur on GBFS v3 feeds, because v3 makes the `is_virtual_station` flag *available* but does not make it *mandatory* or *verified*. The seven-class taxonomy A1–A7 therefore remains operationally relevant under full v3 adoption, and the protocol is straightforwardly extensible to the additional v3 attributes by adding rule-level checks at ingestion time.

3.6 Limitations

The protocol has four acknowledged limitations.

L1 – National coverage. The corpus is limited to metropolitan France; extension to overseas territories and other European jurisdictions requires revisiting thresholds, in particular the macro-regional surface threshold and the geofilter bounds.

L2 – Static versus dynamic. The v1.0 release audits `station_information` only; auditing `station_status` (real-time availability) requires a distinct protocol, sketched in experiment E4 (a 2-day pilot already detected candidate dynamic anomalies on 12.9% of monitored stations). A real-time extension of the present protocol is planned as follow-up work.

L3 – Free-floating static versus dynamic scope. The seven-class taxonomy is applied to free-floating entries as much as to dock-based ones: every free-floating station carries an explicit `station_type = free_floating` label, and the A3 and A7 flags identify the operator-level capacity-publication anomalies that affect them (A3 for over-capacity profiles, A7 for NaN-capacity fields). The published parquet exposes both `capacity_raw` (the raw GBFS value, including placeholders and NaN) and `capacity_audited` (NaN for free-floating entries, dock-based count otherwise), so that downstream consumers have an unambiguous, audit-aware capacity to work with. What this v1.0 release does *not* audit is the *dynamic* behaviour of free-floating fleets – rebalancing patterns, real-time density of available vehicles, per-vehicle relocation – which would require an analogous audit of the `station_status` feed rather than `station_information`. That dynamic audit is out of scope here (it is the topic of L2) ; the present catalogue is therefore best characterised as a *static* audit of GBFS `station_information` across both dock-based and free-floating fleets, with the operational caveat that a free-floating station’s `capacity_audited` value is intentionally null.

L4 – Operator coverage. Some operators publish partial GBFS feeds or are unavailable during the audit window; the inventory completeness is therefore not guaranteed and must be re-evaluated periodically, which is the purpose of the temporal-stability experiment E3.

3.7 Recommendations to operators

Four conventions, compatible with the current GBFS specification, would close most of the anomalies documented here: populate `vehicle_type` exhaustively (A1); leave `capacity` as `null` on virtual stations rather than using placeholders (A2); adopt the GBFS v3 `vehicle_status` extension for free-floating fleets and stop publishing per-vehicle anchors as `station_information` (A3); publish geographical perimeters as polygons in the `geofencing_zones` block (A4, A5); and avoid sentinel coordinates (0,0) on stations whose geolocation is pending, instead omitting them from `station_information` until populated — a single such sentinel in Sevice Seville inflated its bounding-box area by four orders of magnitude in our E5 panel.

Methodological recommendation on A4 / A5 coupling. The Sevice finding (Table 7) suggests that the A5 bounding-box criterion be evaluated on the A4-filtered set of stations, not on the raw set. We retain the as-shipped definition in this paper to make the dependency explicit, but downstream implementations should consider an A5 variant computed after the A4 outlier removal step.

3.8 Future directions

Each of the four limitations above is paired with a falsifiable experiment (Table 11); four of the six include a completed pilot reported in this paper. The remaining priorities are the human inter-rater reliability of E1, the full four-dimensional threshold sweep of E2, the twelve-month temporal stability of E3, and the free-floating-native protocol of E6. Pilot scripts and the final experiment protocols will be packaged under `experiments/` of the source repository alongside the camera-ready release, so that the artefact can be challenged in public rather than defended in private.

3.9 Downstream use

The Audit Catalogue unlocks four families of downstream studies: supply-side composite-index work on cycling-environment quality at station and commune resolution; demand modelling

Table 11: Six falsifiable experiments mapped to the limitations of Section 3.6. “Pilot ✓” indicates a preliminary result already reported in this paper; full protocols will be packaged under `experiments/` of the source repository alongside the camera-ready release.

ID	Title	Targets	Status
E1	Pre-registered held-out audit (post-freeze MobilityData)	Construct validity / circularity	Protocol pre-registered
E2	Threshold and buffer sensitivity	$\sigma_{\max}$, $N_{\min}$, radii	$\sigma_{\max}$ pilot ✓
E3	Twelve-month temporal stability	Snapshot dependence (L4)	Running
E4	Dynamic extension to <code>station_status</code>	L2	2-day pilot ✓
E5	European generalisation (live audit)	L1	10-system panel ✓
E6	Free-floating-native protocol	L3	Planned 2027-Q2

with graph-based neural networks [33] whose false-station nodes from A3 are a known confound; spatial-equity analysis combining station access with INSEE Filosofi; and multimodal integration with GTFS, for which the hybridisation schema is compatible with the canonical GTFS validator [13].

Interactive dashboard. To make the audit verdicts directly inspectable without code, a focused Streamlit dashboard is shipped at <https://gbfs-audit.streamlit.app>, mirrored by `app/streamlit_app.py` in the source repository. The dashboard exposes five tabs: an overview with the corpus-level incidence of A1–A7; a row-level anomaly browser with live multi-flag filters and CSV export; an operator-level audit table with the per-class flag rates that drive the operator-driven hotspots reported in Section 2; a visual data dictionary of the 46 columns with completeness bars; and a free-text data explorer with a KD-tree-backed geographic map coloured by audit confidence. The dashboard is the artefact that turns the audit from a one-off paper into a continuously inspectable reference, addressing the persistence concern raised in Section 3.

4 Methods

4.1 Provenance of the taxonomy

The seven-class taxonomy A1–A7 was built by a hybrid inductive-deductive workflow over six iterations between 2025-04 and 2026-04, starting from the four ISO/IEC 25012 data-quality dimensions (completeness, consistency, accuracy, currentness) and enriched by reverse engineering of both the French free-floating fleets and the global MobilityData catalogue. A1 and A4 emerged from an exploratory scan of the 123 French systems; A5 from enrichment with the *transport.data.gouv.fr* catalogue; A2 from a zero-variance diagnostic on `capacity`; A3 from the per-system median capacity of *Pony* matching a conditional-averaging estimator more closely than a physical-dock count; A6 from the empirical distribution of zero-capacity stations across the 1,509 globally audited systems (Section 4.5); A7 from the bimodal 0%/100% distribution of NaN-capacity stations (Section 4.6). Two rejected hypotheses (“rolling-stock noise”, collapsed into A3; “cross-border duplicate”, no instance found) ship with the audit report for transparency.

4.2 Six-step purging protocol

The transition from the raw feed to the Audit Catalogue proceeds through a chain of six sequential transformations (Algorithm 1), each coded in `utils/data_loader.py` and logged in an audit report shipped alongside the dataset. The protocol is deliberately sequential: early steps are inexpensive and remove out-of-domain or coarse-grained anomalies (A1, A2, A4 perimeter), while later steps require system-level statistics ($\bar{c}$, centroids, σ) and therefore benefit from operating on a pre-cleaned set.

Input: National catalogue $\mathcal{C}$ (123 systems); thresholds $\sigma_{\max} = 3$, $N_{\min} = 20$, $\tau_{A2}^{\text{size}} = 20$, $\tau_{A6} = 0.01$, $\tau_{A7} = 0.5$, perimeter $\Phi = [41^\circ, 52^\circ] \times [-6^\circ, 10^\circ]$.

Output: Certified set $\mathcal{G}$, rejected set $\mathcal{R}$, seven per-row anomaly flags $f_{A1} \dots f_{A7}$, audit report.

$\mathcal{G} \leftarrow \emptyset$; $\mathcal{R} \leftarrow \emptyset$

foreach *system* $s \in \mathcal{C}$ **do**

Ingest GBFS feed $\rightarrow \mathcal{S}_s$

S1. If s is car-sharing, relabel **carsharing** and set $f_{A1} = \text{True}$

S2. If $\text{Var}(c) = 0$, $c > 0$ and $|\mathcal{S}_s^{\text{dock}}| \geq \tau_{A2}^{\text{size}}$, set $f_{A2} = \text{True}$ and force **free_floating**

S3. Compute $\bar{c}_{\text{profile}}/\bar{c}_{\text{actual}}$; if > 5.0 relabel **free_floating** and set $f_{A3} = \text{True}$

S4. Keep $\mathbf{x}_i \in \Phi$; mark stations outside as out-of-perimeter

S5. Compute system bbox area A_s ; if $A_s > 50,000 \text{ km}^2$, set $f_{A5} = \text{True}$ on every $i \in \mathcal{S}_s$

S6. Robust 3- σ outliers on radii $\|\mathbf{x}_i - \boldsymbol{\mu}_s\|$ (median + 1.4826 MAD scaled): set $f_{A4} = \text{True}$

S7. If $\#\{i : c_i = 0 \wedge \text{docked}\}/|\mathcal{S}_s| \geq \tau_{A6}$, set $f_{A6} = \text{True}$ on every $i \in \mathcal{S}_s$

S8. If $\#\{i : c_i = \text{NaN}\}/|\mathcal{S}_s| \geq \tau_{A7}$, set $f_{A7} = \text{True}$ on every $i \in \mathcal{S}_s$

S9. If $|\mathcal{S}_s^{\text{dock}}| < N_{\min}$, label **too_small** and route to $\mathcal{R}$

Update $\mathcal{G}$, $\mathcal{R}$ and the per-row flag matrix accordingly

end

return ($\mathcal{G}$, $\mathcal{R}$, $\{f_{A1} \dots f_{A7}\}$, audit report, metadata)

Algorithm 1: GBFS raw-to-Audit-Catalogue purging pipeline. Steps S1–S3 both purge and label (a station relabelled to **free_floating** or **carsharing** stays in the catalogue with its A-flags set); S4–S5 implement the A5 perimeter check; S6 the A4 outlier check; S7–S8 the A6 / A7 capacity-field anomalies. S9 is the only truly-removing step.

4.3 Protocol properties

Four properties are designed into the pipeline. *Idempotence*: two successive applications produce the same result, tested by a unit test executed on every CI run. *Reversibility*: each excluded station is kept in **rejected_stations.parquet** together with the exclusion motive. *Logging*: every step records the number of stations in, the number of stations out and the type of transformation applied. *Explicit parameterisation*: all thresholds ($\sigma_{\max}$, $N_{\min}$, geofilter bounds, hybridisation radii) are documented as hyperparameters and can be varied for sensitivity analysis without modifying the source code.

These four properties are not arbitrary engineering choices; they operationalise the data-quality dimensions formalised by ISO/IEC 25012 [23] and ISO/IEC 25024 [24], and together they make the pipeline a direct domain-level implementation of the standard. Idempotence enforces the *internal consistency* dimension (two computations on the same input must yield the same output). Reversibility, through the **rejected_stations.parquet** sidecar and the systematic preservation of the exclusion motive, materialises the *recoverability* property and contributes to *traceability*. Logging at every step extends that traceability to the full pipeline run. Explicit hyperparameter declaration addresses the *understandability* dimension at the protocol level. Finally, the FAIR release of the audit report alongside DCAT-AP, JSON Schema and Croissant manifests (Section 4) closes the *compliance* requirement of 25012. Table 12 summarises the mapping; we are not aware of a published reference implementation of ISO/IEC 25012 on an open mobility standard prior to the present work.

Table 12: Mapping of the nine-step purging-protocol design properties to the data-quality dimensions of ISO/IEC 25012:2008 and its companion measurement standard ISO/IEC 25024:2015. The Audit Catalogue is, to our knowledge, the first FAIR-released implementation of these dimensions on a GBFS corpus.

Protocol property	ISO/IEC 25012 dimension	Mechanism in the pipeline
Idempotence	Internal consistency	Unit-tested re-application on certified output
Reversibility	Recoverability + Traceability	<code>rejected_stations.parquet</code> with exclusion motive
Step-level logging	Traceability	Per-step in/out counts and transformation type
Explicit parameterisation	Understandability	Hyperparameters in code + documented in audit report
FAIR release	Compliance + Accessibility	JSON Schema, DCAT-AP, Croissant, ODbL on Zenodo

4.4 Threshold justification

The topological threshold $\sigma_{\max} = 3$ is a compromise between robustness to GPS noise and the preservation of legitimate peripheral extensions, *e.g.* an off-centre TGV station or a university campus. Section 2.5 reports the single-dimension sensitivity pilot: Top-10 invariance for $\sigma_{\max} \geq 2.5$. The robustness threshold $N_{\min} = 20$ dock-based stations is calibrated on the literature of meshed networks: below this value, density and connectivity metrics become noise-dominated and are no longer informative for cross-city comparison [34].

4.5 A6: zero-capacity dock formalisation

The global audit of Section 2.6 surfaced a recurring pattern not covered by A1–A5: a non-trivial fraction of stations in `station_information` declare `capacity = 0`. Because the GBFS schema reserves `station_information` for physical infrastructure and provides `station_status` for transient counts, a non-virtual physical station with zero docks is semantically inconsistent. The signature is therefore:

$$r_s^{A6} = \frac{\#\{i \in \mathcal{S}_s : c_i = 0 \wedge \text{is_virtual_station}(i) \neq \text{true}\}}{|\mathcal{S}_s|} \geq \tau_{A6}. \quad (2)$$

Calibration of τ_{A6} from the empirical distribution. Rather than fixing τ_{A6} by convention, we calibrate it against the population distribution of r_s^{A6} on the 538 globally audited systems with at least 20 stations (Table 13). 94.4 % of audited systems declare *exactly zero* stations with $c = 0$, which establishes the dock-based baseline; the remaining 5.6 % form a heavy tail. From the corresponding empirical CDF we read off two distribution-anchored thresholds. A6-soft uses the 95th percentile $\tau_{A6}^{\text{soft}} = 0.15\%$, which flags any system whose zero-capacity rate is above the corpus median by a non-trivial margin (30 systems globally). A6-hard uses $\bar{r} + 3\hat{\sigma} \approx 2.2\%$, which is the conventional outlier threshold under a normal model on the long tail and corresponds to the empirical 97.5th percentile (8 systems globally). A6 thus exposes a two-tier detection consistent with the rest of the taxonomy: A6-soft as a flag for downstream consumers, A6-hard as a high-confidence operator-side anomaly.

A6 incidence on the audited corpus. Under A6-soft ($\tau = 0.15\%$) 30 of 538 audited systems are flagged; under A6-hard ($\tau = 2.2\%$) 8 systems are flagged. The extreme tail is PBSC HQ (Canada, 18.2 %), Beryl Greater Manchester (UK, 10.5 %), MonaBike (Monaco, 6.8 %), Bicicoruña (Spain, 5.9 %), MBajk (Slovenia, 4.8 %), CyclOcity (Japan, 4.2 %), àVélo (Canada, 3.5 %) and Careem BIKE (UAE, 2.7 %). On the French corpus, no dock-based system crosses either threshold (0/65 dock systems with $c = 0$), which is consistent with the corpus-level median behaviour and provides a negative control. The interaction between A6 and the `is_virtual_station` flag (only required by GBFS v3) is documented as a sensitivity parameter and the rule ships with experiment E1’s inter-rater human validation extended to A6 and A7.

Table 13: Empirical distribution of the zero-capacity rate r_s^{A6} across the 538 globally audited systems with at least 20 stations and a reachable `station_information` feed. Percentiles are computed on the same sample. The two A6 thresholds are derived from the distribution rather than chosen ex ante.

Bucket	Systems	Share
$r_s^{A6} = 0\%$ (no $c = 0$)	508	94.4 %
$0 < r_s^{A6} \leq 0.5\%$	9	1.7 %
$0.5 < r_s^{A6} \leq 1\%$	8	1.5 %
$1 < r_s^{A6} \leq 2\%$	4	0.7 %
$2 < r_s^{A6} \leq 5\%$	6	1.1 %
$5 < r_s^{A6} \leq 10\%$	2	0.4 %
$r_s^{A6} > 10\%$	1	0.2 %
Summary statistics		
mean $\bar{r}$	0.10 %	
standard deviation $\hat{\sigma}$	0.71 %	
95th percentile	0.15 %	
97.5th percentile	$\approx \bar{r} + 3\hat{\sigma} \approx 2.2\%$	
99th percentile	3.54 %	

4.6 A7: null capacity field formalisation

The Italian case study of Section 2.6 demonstrates that A1–A6 are structurally blind to operators that populate `station_information` but leave the `capacity` field unfilled. **A7 is, by design, a consumer-side warning rather than a publisher-side anomaly.** `capacity = NaN` is explicitly permitted by the GBFS v3 specification, in particular for free-floating fleets where the field has no physical referent; a system that publishes `capacity = NaN` on every entry is *spec-compliant*. What A7 captures is the downstream consequence: any analytical pipeline that sums or averages the `capacity` column across such a system aggregates NaN with non-NaN values silently, and produces a number whose physical meaning is undefined. The flag is therefore tooling for the consumer (“do not arithmetically combine these stations with others”), not a verdict on the publisher. The empirical signature is:

$$\rho_s^{A7} = \frac{\#\{i \in \mathcal{S}_s : c_i = \text{NaN}\}}{|\mathcal{S}_s|} \geq \tau_{A7}, \quad (3)$$

with the conservative threshold $\tau_{A7} = 0.5$ chosen so that a system flagged by A7 declares no usable capacity on at least half of its entries. Unlike A6, ρ_s^{A7} is extremely bimodal on the global corpus (Table 14): **87.0 % of audited systems sit at one of the two extreme values 0 % or 100 % NaN**, with only 13.0 % distributed across the interior. Any threshold in $(0, 1)$ partitioning the interior in the same way yields essentially the same flagged set; $\tau_{A7} = 0.5$ is reported for parsimony and is robust to sensitivity sweeps in the tested range.

Incidence of A7. On the 640 globally audited systems with at least 20 stations, **215 systems cross the A7 threshold** and are simultaneously not flagged by any of A1–A6, covering **70,176 stations**. The dominant operator is *Dott* (141 systems across Germany, Italy, Austria, UAE, and other markets), followed by *nextbike* (13 systems where the operator chose NaN rather than placeholder), *Bird* (10 Canadian deployments) and several smaller operators. The empirical conclusion is that the same publication choice that affects the French Dott and Bird fleets is propagated internationally, and that the v1.0 taxonomy A1–A5 captures only the visible failure modes of the field; A6 and A7 jointly close the principal blind spots that remain on the `capacity` channel.

Table 14: Empirical distribution of the NaN-capacity rate ρ_s^{A7} across the 640 globally audited systems with at least 20 stations and a reachable `station_information` feed. The corpus is sharply bimodal at the extreme values, which justifies a coarse threshold rather than a distribution-anchored one.

Bucket	Systems	Share
$\rho_s^{A7} = 0\%$ (no NaN)	238	37.2 %
$0 < \rho_s^{A7} < 1\%$	4	0.6 %
$1 \leq \rho_s^{A7} < 10\%$	11	1.7 %
$10 \leq \rho_s^{A7} < 50\%$	22	3.4 %
$50 \leq \rho_s^{A7} < 99\%$	41	6.4 %
$99 \leq \rho_s^{A7} < 100\%$	5	0.8 %
$\rho_s^{A7} = 100\%$ (all NaN)	319	49.8 %
Mass at the two extremes	557	87.0 %

4.7 Multi-source hybridisation

The output of the purging stage enriches each station with a contextual profile obtained by articulating six reference sources (Table 15). Hybridisation is performed either by spatial joining (point-in-polygon, buffer queries) or by administrative aggregation (*i.e.* at the IRIS or municipal level), depending on the granularity of the source. Each row of Table 15 corresponds to one or more typed columns of the released parquet schema (Table 16).

Table 15: Reference sources used for contextual enrichment. Every listed source contributes at least one column to the released parquet schema (Table 16). Out-of-scope sources considered during prototyping but not retained in the v1.0 release (OpenStreetMap point-of-interest density, FUB cycling barometer score) are documented in the audit report.

Source	Granularity	Variables	Operation
INSEE Filosofi	IRIS / municipality	Median income/CU, local Gini, decile, car-less share	Aggregation
INSEE Recensement	municipality	Share of commute by bike (<code>part_velo_travail</code>)	Aggregation
IGN BD ALTI	DEM raster	Elevation, topographic roughness index	Spatial samp
IGN BD TOPO	polylines	Cycle-infrastructure linear (300 m buffer)	Spatial join
ONISR BAAC (Cerema)	geolocated accident	Severe-crash count (500 m, 5 yr)	Spatial join
National GTFS aggregator	GTFS stops	Heavy-transit stops within 300 m	Spatial join

4.8 Released schema

The Audit Catalogue Parquet exposes 46 columns documented in Table 16: 14 from the raw GBFS payload (preserved verbatim or normalised), 11 from the audit pipeline (typed `station_type` enum, capacity-raw and capacity-audited pair, seven per-station anomaly flags A1–A7, operator label, and audit-confidence rating), 5 from intra- and inter-system network geometry, 4 from spatial-equity hybridisation, 3 from infrastructure hybridisation, and 2 from multimodal-access hybridisation. Every column carries a stable name, a declared dtype, a source pointer and a completeness rate; the machine-readable equivalent is shipped as a JSON Schema, a Frictionless Data Package descriptor and a Croissant manifest. The eleven audit-pipeline columns are the key novelty of the catalogue: they make the audit’s verdict inspectable per row, rather than aggregated in a separate report.

Empirical completeness on the 5,442 dock-based stations is above 93 % for every contextual variable except the optional `address` field of the raw GBFS payload (Figure 7). The minimum-completeness threshold of 90 % documented in the audit report is therefore reached on every certified column of the released schema.




Figure 7: Empirical completeness of contextual enrichment per derived variable on the dock-based subset of the Audit Catalogue ($n = 5,442$). All variables exceed the 90 % completeness threshold required by the audit report; missing values remain below 6 % across the corpus.

4.9 Reproducibility infrastructure

The Audit Catalogue is released in alignment with the four FAIR principles [35]. *Findable*: primary deposit on Zenodo under concept DOI 10.5281/zenodo.20125460, with one DOI per SemVer-tagged release (current 1.0.0, dated 2026-05); mirrored on *Recherche Data Gouv* (HAL/BSO indexing). *Accessible*: Parquet format, direct HTTPS download and Git LFS mirror. *Interoperable*: the schema extends GBFS v3 and is published as a JSON Schema, a Frictionless Data Package descriptor [17], a DCAT-AP record [36] and a Croissant manifest [37]. *Reusable*: released under ODbL v1.0; code under MIT, with audit report and source code. Raw sources are archived with SHA-256 hashes; the Python environment is pinned in `requirements.txt`; a Docker image is published for bit-exact reproduction. The audit pipeline itself is fully deterministic (no random sampling); the only stochastic step is the bootstrap resampling used for the confidence intervals reported in the main text, for which the random seed is pinned at `numpy.random.default_rng(42)`. Idempotence and per-class detector unit tests run in CI on every push; the audit report is regenerated automatically on every release tag.

4.10 Implementation

The released Parquet file is 7.8 MB on disk for the 46,307 stations and loads in under one second into a standard `pandas` session on commodity hardware ; typical analytical queries (city-level filtering, bounding-box selection, joins with external panels) complete in tens of milliseconds. The audit pipeline is implemented as a compact Python package (`audit_pipeline`, ~ 560 lines of code across three modules: `core` for the Tier-1 / Tier-2 detectors, `regenerate` for one-command artefact refresh, and the package `__init__` for the public API) with stable function-level entry points documented in the source repository.

Software validation. The audit primitives are accompanied by a unit-test suite (`tests/test_core.py`, 24 tests, 85 % line coverage of the public `audit_pipeline.core` surface) that exercises each of the seven anomaly classes in isolation as well as the combined end-to-end `enrich()` pipeline. Each detector is paired with a deterministic synthetic fixture engineered to trip exactly that class (e.g., a 25-station grid with one transposed-coordinate row for A4, a 32-station fleet spanning more than 300,000 km² for A5, a 60 %-NaN-capacity dock-based system for A7) and with specificity controls verifying that healthy fixtures stay unflagged. The suite further checks that the Tier-2 geometry columns (intra- and inter-system k -nearest-neighbour distances, local densities) are consistent with the equirectangular projection used for the meter-scale arithmetic, and that the pipeline is idempotent under repeated application. Tests are executed by `pytest` as a GitHub Actions workflow on every push and pull request, and locally inside the published Docker image.

We stress that this suite validates the correspondence between the implemented logic and the published anomaly taxonomy; it is not a substitute for the hand-curated cross-operator audit, which provides the external ground truth against which the pipeline’s verdicts were calibrated.

5 Data and code availability

Lead contact. Requests for further information should be directed to and will be fulfilled by the lead contact, Rohan Fossé (rfosse@cesi.fr).

Materials availability. This study did not generate new physical materials. The released artefact is the *GBFS France Audit Catalogue* dataset, a derivative of public open data.

Data and code. The artefacts cover two scopes with explicit licensing boundaries.

- **French Audit Catalogue (Zenodo, ODbL v1.0).** The deposit contains four artefacts under the same concept DOI 10.5281/zenodo.20125460 (current release 1.0.0, 2026-05), mirrored on *Recherche Data Gouv*: (i) the 46,307-station certified Parquet (`stations_gold_standard_final.parquet`, ~ 6.6 MB, 46 typed columns including per-row anomaly flags and audited capacities) ; (ii) the lightweight system-level executive summary (`stations_gold_standard_audit_summary.csv`, ~60 kB, one row per audited system with anomaly counts and operator attribution) for users who want an overview before downloading the full parquet ; (iii) the `rejected_stations.parquet` log with exclusion motives ; and (iv) the audit report and machine-readable descriptors (JSON Schema, DCAT-AP, Frictionless Data Package, Croissant JSON-LD). The 142 candidate French GBFS feeds inventoried from *transport.data.gouv.fr* and MobilityData are archived with SHA-256 hashes in the same deposit ; the certified subset of 123 systems is derived deterministically from them.
- **Source code (MIT licence).** The audit pipeline (`audit_pipeline` Python package, nine-step purging protocol, hybridisation joins, FAIR exporters) is released at <https://github.com/rohanfosse/gbfs-audit-catalogue>, together with the manuscript source, the companion Jupyter notebook (`notebooks/catalogue_recipes.ipynb`) and a focused interactive dashboard with five tabs (Overview, Anomaly browser, Operator audit, Schema, Data explorer) at `app/streamlit_app.py`, deployed publicly at <https://gbfs-audit.streamlit.app>. A CITATION.cff file and a Docker image (`ghcr.io/rohanfosse/gbfs-audit:1.0.0`) are provided for bit-exact reproduction. Releases are mirrored to Zenodo via the GitHub–Zenodo integration.
- **Hugging Face mirror.** The released parquet is mirrored as a Hugging Face dataset at <https://huggingface.co/datasets/rohanfosse/gbfs-audit-catalogue> for one-line access via `datasets.load_dataset()`.
- **Global audit reproduction.** The full per-system result of the 1,509-system MobilityData-catalogue audit is versioned as code rather than as data, because the upstream catalogue evolves. A re-run is one command: `python -m audit_pipeline.global_audit` from the source repository. The output (`massive_audit_results.csv` and `massive_audit_summary.json`) is re-generated on every release and archived as a Zenodo asset attached to the corresponding tag.
- **Additional resources.** The 500-station blinded sample used in the internal robustness checks, the orthogonal- and full-task classifier scripts, the $\sigma_{\max}$ sensitivity notebook, and the per-operator characterisation of the French free-floating subset are released as supplementary material alongside the dataset.

A Catalogue use cases

Five concrete reuse patterns are documented below. Each is a self-contained few-line snippet against `stations_gold_standard_final.parquet` ; a complete Jupyter notebook (`catalogue_recipes.ipynb`) is released alongside the dataset.

Use case 1 – Loading the catalogue. A single `read_parquet` call exposes the 46 typed columns ; typical project filters use `station_type`, `audit_confidence` and the per-row anomaly flags `flag_A1..flag_A7`.

```
import pandas as pd
gs = pd.read_parquet("stations_gold_standard_final.parquet")
clean = gs[(gs.station_type == "docked_bike") &
            (gs.audit_confidence == "high")]
len(clean) # returns 5,402
```

Use case 2 – Per-operator anomaly profile. The audit's verdict is exposed at the row level by the seven flag columns, so a single `groupby` returns the per-operator anomaly rates that drive Section 2 :

```
gs.groupby("operator_name").agg(
    n=("uid", "size"),
    A3_rate=("flag_A3", "mean"),
    A7_rate=("flag_A7", "mean"),
).sort_values("n", ascending=False).head(10)
```

The two largest operators on the corpus (*Dott*, $\sim 20\,000$ stations ; *Pony*, $\sim 11\,500$) come out at `A7_rate = 1.0` and `A3_rate = 1.0` respectively, which is the row-level expression of the operator-driven hotspots documented in Section 2.

Use case 3 – Raw vs. audited capacity per system. `capacity_raw` preserves the GBFS-declared value (including placeholders and NaN) while `capacity_audited` carries the post-audit count (NaN for non-dock types) ; one can therefore quantify per-system the gap between what GBFS *claims* and what the audit *certifies* :

```
sys = gs.groupby("system_id").agg(
    raw_capacity=("capacity_raw", "sum"),
    audited_capacity=("capacity_audited", "sum"),
).query("raw_capacity > 1000")
sys["over_count_factor"] = sys.raw_capacity /
sys.audited_capacity.replace(0, 1)
sys.sort_values("over_count_factor", ascending=False).head(5)
```

Use case 4 – Spatial-equity correlation. Per-city correlation between dock-based station count and median income per consumption unit :

```
by_city = (
    gs[gs.station_type == "docked_bike"]
    .groupby("city").agg(
        n=("uid", "size"),
        income=("revenu_median_uc", "median"))
)
rho = by_city.n.corr(by_city.income, method="spearman")
# rho  $\approx +0.09$  on 64 cities (weak)
```

The weak positive coupling at the city level is consistent with the much stronger structural pattern uncovered by use case 5 below : the income-cycling gap is not a city-mean effect but a within-city station-level effect concentrated in the lower income quartile.

Use case 5 – Identifying mobility deserts. Three filters on `station_type`, `revenu_median_uc` and the multimodal `gtfs_heavy_stops_300m` count materialise the desert concept of Section 2 :

```
q1 = gs.revenu_median_uc.quantile(0.25)
deserts = gs[(
    gs.station_type == "docked_bike") &
    (gs.revenu_median_uc < q1) &
    (gs.gtfs_heavy_stops_300m == 0)]
deserts.city.value_counts().head(10)
```

This three-condition filter recovers 1,682 stations (30.9 % of the dock-based corpus) with a Q1 income threshold of 21,040 EUR per consumption unit on the v1.0 release of the catalogue. Figure 8 shows the top fifteen contributing cities; the count is sensitive to the income data vintage and to the GTFS-stop buffer, both of which are documented hyperparameters of the audit.




Figure 8: Top fifteen French urban areas by number of dock-based mobility-desert stations, recovered by the three-condition filter of Example 3. A station counts as a mobility desert if its commune sits in the lowest national income quartile and the station has zero heavy-transit stops within a 300 m buffer.

Acknowledgments

The authors thank the producers of the open data that made this research possible: GBFS operators, MobilityData, OpenStreetMap, Cerema, ONISR, INSEE, FUB and the French GTFS community.

CRedit authorship contribution statement

Rohan Fossé: Conceptualization, Methodology, Software, Data curation, Formal analysis, Validation, Visualization, Writing – original draft, Writing – review & editing. **Gaël Pallares:** Conceptualization, Methodology, Supervision, Project administration, Writing – review & editing.

Declaration of competing interests

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Declaration of generative AI and AI-assisted technologies in the writing process

During the preparation of this work the authors used GitHub Copilot and Claude (Anthropic) for code completion and editorial language polishing. After using these tools the authors reviewed and edited the content as needed and take full responsibility for the content of the publication.

Funding

This work received no specific grant from any funding agency.

References

- [1] John Pucher, Jennifer Dill, and Susan Handy. Infrastructure, programs, and policies to increase bicycling: An international review. *Preventive Medicine*, 50:S106–S125, 2010.
- [2] MobilityData. General bikeshare feed specification (gbfs) version 3.0. <https://github.com/MobilityData/gbfs>, 2024. Open standard, Open Mobility Foundation.
- [3] Open Mobility Foundation. Gbfs – open standard for real-time information of shared mobility systems. <https://www.openmobilityfoundation.org/about-gbfs/>, 2023.
- [4] Ralph Buehler and John Pucher. Trends in walking and cycling safety: Recent evidence from high-income countries, with a focus on the United States and Germany. *American Journal of Public Health*, 107(2):281–287, 2017.
- [5] Cyrille Médard de Chardon, Geoffrey Caruso, and Isabelle Thomas. Bike-share rebalancing strategies, patterns, and purpose. *Journal of Transport Geography*, 64:218–233, 2017.
- [6] Ezgi Eren and Volkan Emre Uz. A review on bike-sharing: The factors affecting bike-sharing demand. *Sustainable Cities and Society*, 54:101882, 2020.
- [7] Marie-Pier Pelletier, Martin Trépanier, and Catherine Morency. Smart card data use in public transit: A literature review. *Transportation Research Part C: Emerging Technologies*, 19(4):557–568, 2011.
- [8] Sui Tao, David Rohde, and Jonathan Corcoran. Examining the spatial-temporal dynamics of bus passenger travel behaviour using smart card data and the flow-comap. *Journal of Transport Geography*, 41:21–36, 2020.
- [9] Gustavo Romanillos, Martin Zaltz Austwick, Dick Ettema, and Joost De Kruijf. Big data and cycling. *Transport Reviews*, 36(1):114–133, 2016.
- [10] Julie Bachand-Marleau, Brian H. Y. Lee, and Ahmed M. El-Geneidy. Better understanding of factors influencing likelihood of using shared bicycle systems and frequency of use. *Transportation Research Record*, 2314:66–71, 2012.
- [11] Elliot Fishman. Bikeshare: A review of recent literature. *Transport Reviews*, 36(1):92–113, 2016.

- [12] North American Bikeshare and Scootershare Association. 2023 State of the Industry Report. <https://nabsa.net/about/industry/>, 2024. Annual cross-operator data harmonisation report.
- [13] MobilityData. Canonical GTFS schedule validator. <https://github.com/MobilityData/gtfs-validator>, 2024. Reference implementation of the GTFS schedule validation rules.
- [14] Aaron Antrim and Sean J. Barbeau. The many uses of GTFS data – opening the door to transit and multimodal applications. In *ITS America 23rd Annual Meeting Conference Proceedings*, Nashville, TN, 2013.
- [15] Great Expectations. Great expectations: a python framework for declarative data validation, documentation and profiling. <https://greatexpectations.io>, 2024.
- [16] Niels Bantilan. Pandera: Statistical data validation for pandas. <https://pandera.readthedocs.io>, 2024.
- [17] Open Knowledge Foundation. Frictionless data specifications: Data package. <https://specs.frictionlessdata.io/data-package/>, 2023. Open standard for describing and packaging tabular data.
- [18] Richard Y. Wang and Diane M. Strong. Beyond accuracy: What data quality means to data consumers. *Journal of Management Information Systems*, 12(4):5–33, 1996.
- [19] Leo L. Pipino, Yang W. Lee, and Richard Y. Wang. Data quality assessment. *Communications of the ACM*, 45(4):211–218, 2002.
- [20] Vijay Khatri and Carol V. Brown. Designing data governance. *Communications of the ACM*, 53(1):148–152, 2010.
- [21] Laura Sebastian-Coleman. *Measuring Data Quality for Ongoing Improvement: A Data Quality Assessment Framework*. Morgan Kaufmann, Waltham, MA, 2013.
- [22] Felix Naumann. Data profiling revisited. *ACM SIGMOD Record*, 42(4):40–49, 2014.
- [23] ISO/IEC. *ISO/IEC 25012:2008 – Software engineering – Software product Quality Requirements and Evaluation (SQuaRE) – Data quality model*. International Organization for Standardization, Geneva, 2008.
- [24] International Organization for Standardization. ISO/IEC 25024:2015: Systems and software engineering — systems and software quality requirements and evaluation (SQuaRE) — measurement of data quality. Standard, Geneva, Switzerland, 2015.
- [25] International Organization for Standardization. ISO 19157:2013: Geographic information — data quality. Standard, Geneva, Switzerland, 2013.
- [26] European Parliament and Council. Directive 2007/2/ec establishing an infrastructure for spatial information in the european community (INSPIRE). Official Journal of the European Union L 108/1, 2007.
- [27] World Wide Web Consortium. Data catalog vocabulary (DCAT) — version 2. W3C Recommendation, <https://www.w3.org/TR/vocab-dcat-2/>, 2020.
- [28] World Wide Web Consortium. Data on the Web best practices: Data quality vocabulary. W3C Working Group Note, <https://www.w3.org/TR/vocab-dqv/>, 2016.

- [29] Marijn Janssen, Yannis Charalabidis, and Anneke Zuiderwijk. Benefits, adoption barriers and myths of open data and open government. *Information Systems Management*, 29(4):258–268, 2012.
- [30] Susan Leigh Star. This is not a boundary object: Reflections on the origin of a concept. *Science, Technology, & Human Values*, 35(5):601–617, 2010.
- [31] Michael Stonebraker and Uğur Çetintemel. “one size fits all”: An idea whose time has come and gone. In *Proceedings of the 21st International Conference on Data Engineering (ICDE)*, pages 2–11, 2005.
- [32] John Hughes. QuickCheck testing for fun and profit. *Practical Aspects of Declarative Languages (PADL)*, pages 1–32, 2007.
- [33] Lei Lin, Zhengbing He, and Srinivas Peeta. Predicting station-level hourly demand in a large-scale bike-sharing network: A graph convolutional neural network approach. *Transportation Research Part C*, 97:258–276, 2018.
- [34] Meead Saberi, Mehrnaz Ghamami, Yi Gu, Mohammad Hassan Saadi Shojaei, and Elliot Fishman. Understanding the impacts of a public transit disruption on bicycle sharing mobility patterns: A combined graph and hyper-graph approach. *Transportmetrica A: Transport Science*, 14(7):576–600, 2018.
- [35] Mark D. Wilkinson, Michel Dumontier, IJsbrand Jan Aalbersberg, Gabrielle Appleton, Myles Axton, Arie Baak, et al. The FAIR guiding principles for scientific data management and stewardship. *Scientific Data*, 3(160018), 2016.
- [36] European Commission. DCAT-AP: DCAT application profile for data portals in europe, version 2.1.0. <https://semiceu.github.io/DCAT-AP/>, 2021.
- [37] Mubashara Akhtar, Omar Benjelloun, Costanza Conforti, Luca Foschini, et al. Croissant: A metadata format for ML-ready datasets. <https://github.com/mlcommons/croissant>, 2024. MLCommons Croissant working group, JSON-LD schema adopted by Hugging Face, Kaggle and Google Dataset Search.

Table 16: Released schema of `stations_gold_standard_final.parquet`. “Compl.” is the empirical completeness rate on the certified subset. Columns marked [†] are introduced by the audit pipeline; columns marked * are derived from external sources.

Column	Type	Source / Description	Compl.
<i>Identifiers (5)</i>			
<code>uid</code>	string	Audit Catalogue primary key	100 %
<code>station_id</code>	string	GBFS native identifier	100 %
<code>system_id</code>	string	Operator-system identifier	100 %
<code>system_name</code>	string	Operator-system label	100 %
<code>source</code>	string	Feed URL / catalogue source	100 %
<i>Spatial and administrative (5)</i>			
<code>lat, lon</code>	float64	Geofiltered WGS84 coordinates	100 %
<code>city</code>	string	Normalised city name	100 %
<code>commune_name</code>	string	INSEE commune label	99.1 %
<code>code_commune</code>	string	INSEE commune code	99.1 %
<code>region_id</code>	string	Administrative region	100 %
<i>Station description (4)</i>			
<code>station_name</code>	string	GBFS station name	99.8 %
<code>address</code>	string	GBFS address	41.2 %
<code>capacity</code>	int	Declared capacity (may be a placeholder)	96.8 %
<code>n_stations_system</code>	int	Total stations in parent system	100 %
<i>Audit pipeline outputs (11)[†]</i>			
<code>station_type</code>	enum	{docked_bike, free_floating, carsharing}	100 %
<code>capacity_raw</code>	float64	Raw GBFS capacity (may be NaN or placeholder)	30.1 %
<code>capacity_audited</code>	float64	Audited dock count (NaN for non-dock types)	30.1 %
<code>flag_A1</code>	bool	Out-of-domain inclusion (carsharing)	100 %
<code>flag_A2</code>	bool	Placeholder capacity at system level	100 %
<code>flag_A3</code>	bool	Structural over-capacity (free-floating)	100 %
<code>flag_A4</code>	bool	Geospatial outlier (3- σ on system centroid)	100 %
<code>flag_A5</code>	bool	Out-of-perimeter (system bbox > 50,000 km ²)	100 %
<code>flag_A6</code>	bool	Zero-capacity dock	100 %
<code>flag_A7</code>	bool	Null-capacity field at system level	100 %
<code>operator_name</code>	string	Normalised operator label	100 %
<code>audit_confidence</code>	enum	{high, medium, low}	100 %
<code>fetch_at</code>	datetime	Timestamp of the audited snapshot	100 %
<i>Network geometry (5)[†]</i>			
<code>dist_to_nearest_station_m</code>	float64	Intra-system KNN distance	99.6 %
<code>n_stations_within_500m</code>	int	Intra-system 500 m density	100 %
<code>n_stations_within_1km</code>	int	Intra-system 1 km density	100 %
<code>nearest_system_dist_m</code>	float64	Distance to any other system	100 %
<code>catchment_density_per_km2</code>	float64	Stations per km ² (1 km buffer)	100 %
<i>Topography (2)*</i>			
<code>elevation_m</code>	float64	BD ALTI elevation	99.4 %
<code>topography_roughness_index</code>	float64	Local relief amplitude	98.1 %
<i>Cycling infrastructure (2)*</i>			
<code>infra_cyclable_km</code>	float64	BD TOPO cycle-lane linear (300 m)	97.6 %
<code>infra_cyclable_pct</code>	float64	Share of dedicated right-of-way	96.9 %
<i>Safety (1)*</i>			
<code>baac_accidents_cyclistes</code>	int	Severe-crash count (500 m, 5 y)	99.8 %
<i>Multimodal access (2)*</i>			
<code>gtfs_heavy_stops_300m</code>	int	Heavy-transit stops within 300 m	98.7 %
<code>gtfs_stops_within_300m_pct</code>	float64	Share of accessible heavy-transit	98.7 %
<i>Socio-economy (5)*</i>			
<code>revenu_median_uc</code>	float64	INSEE Filosofi median income/CU	95.4 %
<code>gini_revenu</code>	float64	Log30 Gini index	93.0 %
<code>revenu_d1</code>	float64	First-decile income	95.2 %
<code>ecart_interquar</code>	float64	Interquartile income spread	95.2 %
<code>part_menages_voit0</code>	float64	Share of car-less households	95.4 %